



Matthew Fili <matthewjosephfili@gmail.com>

Anthropic IPO filing 📄, OpenAI on AWS ☁️, Perplexity search code 🔍

1 message

TLDR AI <dan@tldrnewsletter.com>

Tue, Jun 2, 2026 at 8:49 AM

To: matthewjosephfili@gmail.com

[Sign Up](#) | [Advertise](#) | [View Online](#)

TLDR

Together With  Dataiku**TLDR AI 2026-06-02****[Dataiku drives the shift to always-on AI governance \(Sponsor\)](#)**

By integrating governance directly into AI workflows, [Dataiku supports enterprise-scale AI](#) through:

- Embedded controls across the AI lifecycle
- Open, model-agnostic architecture
- Continuous oversight without operational drag

[In this exclusive excerpt](#), learn how IDC MarketScape assesses AI governance platforms, and what CIOs and data leaders should prioritize when governing [analytics](#), [ML](#), [GenAI](#), and [AI agents at scale](#).

**Headlines & Launches****[Anthropic Filed a Confidential Draft IPO Registration \(2 minute read\)](#)**

Anthropic confidentially submitted a draft S-1 registration statement to the US Securities and Exchange Commission for a proposed initial public offering. The filing does not set pricing or share counts and remains subject to regulatory review and market conditions.

[OpenAI and Codex Reach AWS \(3 minute read\)](#)

OpenAI announced the general availability of its frontier models and Codex on AWS. The integration allows enterprises to access OpenAI capabilities through existing AWS security, governance, procurement, and billing workflows.

[NVIDIA just announced the release of Nemotron 3 Ultra \(2 minute read\)](#)

Nemotron 3 Ultra features 550B parameters (55B active). It is the most intelligent open weights model from the US. The model will be made available in NVFP4 quantization for higher inference performance. It scores 48 on the Artificial Analysis Intelligence Index, well ahead of the next strongest model, Gemma 4 31B, which scored 39. Nemotron 3 Ultra serves over 300 tokens per second on a pre-release Deep Infra endpoint.

[Qwen3.7-Plus: Multimodal Agent Intelligence \(36 minute read\)](#)

Qwen3.7-Plus is a multimodal agent model that unifies vision and language into a single, versatile agent foundation. It can operate as a multimodal interactive hybrid agent, seamlessly blending GUI and CLI interactions within a single agent loop. The model performs consistently across scaffolds and frameworks. Qwen3.7-Plus is now available via Alibaba Cloud Model Studio.



Deep Dives & Analysis

[Why Video Agent models are next — Ethan He, xAI Grok Imagine \(98 minute read\)](#)

Ethan He was the lead on Nvidia's Cosmos World Model. He then joined xAI and built Grok Image in three months. He has been in the center of some of the most important work in video generation, multimodal models, and real-time world models. This post contains an interview with He where he unpacks what it actually takes to build frontier image and video systems.

[Opus 4.8 Part 2: Model Welfare \(42 minute read\)](#)

Anthropic cares about model welfare, and it attempts to address it through studies. The study of model welfare is difficult, and Anthropic mainly relies on self-reporting by the model. It can be difficult to evaluate whether model responses are really representative of the truth. This article takes a look at Anthropic's findings regarding model welfare with Opus 4.8.



Engineering & Research

[1,000+ Datadog customers use AI in prod. Here's what the LLM telemetry shows \(Sponsor\)](#)

Ever wonder how your competitors are actually using AI? Datadog's [State of AI Engineering](#) report uses LLM telemetry from over 1,000 orgs to show you how model provider adoption is changing, why LLM tech debt is already compounding, and where those hidden token costs are coming from. [Get your copy](#)

[Running OpenAI Models on Amazon Bedrock \(58 minute read\)](#)

OpenAI cookbook walks through building production workflows with OpenAI models hosted on Amazon Bedrock using the Responses API. It covers structured outputs, tool calling, file inputs, state management, prompt caching, and operational best practices.

[Rethinking Search as Code Generation \(25 minute read\)](#)

Perplexity introduces Search as Code (SaC) to modernize search architectures by allowing

models direct control over the search process via an SDK. This approach lets AI models configure search pipelines tailored to specific tasks, improving performance and efficiency over traditional monolithic systems. SaC outperformed competitors in benchmarks, especially in complex tasks like WANDR, demonstrating its robust, cost-effective agentic search capabilities.

[NVIDIA Launches Cosmos 3, the Open Frontier Foundation Model for Physical AI \(5 minute read\)](#)

Nvidia Cosmos 3 is a new leaderboard-topping open physical AI foundation model. It is a fully open omnimodel with native vision reasoning and multimodal generation across text, image, video, ambient sound, and action. The model is built on a mixture-of-transformer architecture, which pairs a reasoning transformer with an expert generation transformer. It gives developers a powerful pretrained foundation for building physical AI systems with less data and lower training costs.

[JetBrains's Mellum 2 \(49 minute read\)](#)

JetBrains introduced Mellum 2, a 12B-parameter MoE language model optimized for coding, reasoning, tool use, and agentic workflows.



Miscellaneous

[US moves to close the loophole letting Nvidia's top chips reach Chinese firms abroad \(3 minute read\)](#)

The US Commerce Department has issued guidance extending export license requirements to advance chips sold to any entity headquartered in China, regardless of where that entity is physically located. This closes a loophole that allowed Chinese companies to buy Nvidia chips using a subsidiary in another country. The action is aimed at future sales and not at clawing back hardware already shipped. It doesn't affect the servicing of advanced computing equipment such as servers.

[Alphabet plans to raise \\$80 billion from stock sales to fund AI buildout \(4 minute read\)](#)

Alphabet plans to sell \$80 billion in stock to fund investments into AI compute infrastructure due to unprecedented customer demand. The raise includes \$10 billion from Berkshire Hathaway, \$30 billion in an underwritten offering, and \$40 billion from an at-the-market offering program for Class A and Class C shares expected to begin in the third quarter. Goldman Sachs, JPMorgan Chase, and Morgan Stanley are acting as joint book-running managers for the underwritten offering.



Quick Links

TLDR is hiring a Senior Software Engineer, Applied AI (\$250k-\$350k, Fully Remote)

TLDR's Applied AI team is tasked with making every process at TLDR legible to code, runnable by anyone, and composable into larger workflows. Join a small, fast moving team using the latest AI tools with an unlimited token budget. [Learn more](#).

Opus 4.8 just broke ARC-AGI-3 (1 minute read)

It tripled GPT-5.5's score.

Mistral Search Toolkit for Production AI Pipelines (4 minute read)

Mistral released Search Toolkit in public preview, an open-source framework that unifies data ingestion, retrieval, and evaluation within a shared interface.

Cursor Expands Teams Usage Limits (2 minute read)

Cursor announced higher Teams plan limits, a new Premium seat for heavy agent users, and additional spending controls for administrators.

Love TLDR? Tell your friends and get rewards!

Share your referral link below with friends to get free TLDR swag!

<https://refer.tldr.tech/58d35097/2>

[Track your referrals here.](#)

Want to advertise in TLDR?

If your company is interested in reaching an audience of AI professionals and decision makers, you may want to [advertise with us](#).

Want to work at TLDR?

[Apply here](#), [create your own role](#) or send a friend's resume to jobs@tldr.tech and get \$1k if we hire them! TLDR is one of [Inc.'s Best Bootstrapped businesses](#) of 2025.

If you have any comments or feedback, just respond to this email!

Thanks for reading,
[Andrew Tan](#), [Ali Aminian](#), & [Jacob Turner](#)

[Manage your subscriptions](#) to our other newsletters on tech, startups, and programming. Or if TLDR AI isn't for you, please [unsubscribe](#).